



Optimization in Machine Learning

First order methods

Step size and optimality




Learning goals

- Impact of step size
- Fixed vs. adaptive step size
- Exact line search
- Armijo rule & Backtracking
- Bracketing & Pinpointing

Slides based on [Wright and Recht 2022](#) (Ch. 3.5)



OVERVIEW

- Iteration t : Choose not only descent direction $\mathbf{d}^{[t]}$, but also step size $\alpha^{[t]}$
- Three main approaches:
 - ❶ **Fixed** step size: constant α
 - ❷ **Adaptive** step size: $\alpha^{[t]}$ is adjusted over time
 - ❸ **Line search**: choose $\alpha^{[t]}$ by (approximately) minimizing f along $\mathbf{d}^{[t]}$
 - *Exact*: find optimal $\alpha^{[t]}$
 - *Inexact*: find “good enough” $\alpha^{[t]}$




FIXED STEP SIZE

- Fixed step size: $\alpha^{[t]} = \alpha > 0$ constant across all iterations
- **Difficulty:** How to choose α properly?
 - If α too small, procedure may converge very slowly (left)
 - If α too large, procedure may not converge $\rightarrow$ “jumps” around optimum (right)




Steps of GD for $f(\mathbf{x}) = 10x_1^2 + x_2^2/2$

- Approaches: (i) trial and error, or (ii) choosing α based on global properties of f (e.g., Lipschitz constant L of ∇f , next slide deck)



ADAPTIVE STEP SIZE

- How can we adaptively control step size?
- A natural way is to decrease $\alpha^{[t]}$ over time, e.g., $\alpha^{[t]} = \frac{10}{t}$
- Many different adaption policies exist ▶ Zhang et al. 2023
- **Example:** GD on $f(x) = \begin{cases} \frac{1}{2}x^2 & \text{if } |x| \leq \delta, \\ \delta \cdot (|x| - 1/2 \cdot \delta) & \text{otherwise.} \end{cases}$



GD with small constant (**red**), large constant (**green**), and diminishing (**blue**) step size

— $\alpha = 10$




— $\alpha = 10/t$

EXACT LINE SEARCH



- Use **optimal** step size in each iteration:

$$\alpha^{[t]} = \arg \min_{\alpha \in \mathbb{R}_{>0}} g(\alpha) = \arg \min_{\alpha \in \mathbb{R}_{>0}} f(\mathbf{x}^{[t]} + \alpha \mathbf{d}^{[t]})$$

- Need to solve a **univariate** optimization problem in each iteration $\Rightarrow$ univariate optimization methods
- **Problem:** Expensive, prone to poorly conditioned problems
- **But:** No need for *optimal* step size. Only need a step size that is “good enough”.
- **Reason:** Effort may not pay off, but in some cases slows down performance.





ARMIJO RULE



- **Inexact line search:** Minimize objective “sufficiently” without computing optimal step size exactly
- Common condition to guarantee “sufficient” decrease: **Armijo rule**

ARMIJO RULE





- Fix $\gamma_1 \in (0, 1)$. α satisfies **Armijo rule** in $\mathbf{x}$ for descent direction $\mathbf{d}$ if

$$f(\mathbf{x} + \alpha \mathbf{d}) \leq f(\mathbf{x}) + \gamma_1 \alpha \nabla f(\mathbf{x})^T \mathbf{d}.$$

- Note:** $\nabla f(\mathbf{x})^T \mathbf{d} < 0$ ($\mathbf{d}$ descent dir.) $\implies f(\mathbf{x} + \alpha \mathbf{d}) < f(\mathbf{x})$.

ARMIJO RULE





- **Feasibility:** For descent direction $\mathbf{d}$ and $\gamma_1 \in (0, 1)$, there exists $\alpha > 0$ fulfilling Armijo rule. In many cases, Armijo rule guarantees local convergence of GD and is therefore frequently used.

BACKTRACKING LINE SEARCH

- Procedure to meet the Armijo rule: **Backtracking** line search
- **Idea:** Decrease α until Armijo rule is met

Algorithm Backtracking line search

- 1: Choose initial step size $\alpha = \alpha_{\text{init}}$, $0 < \gamma_1 < 1$ and $0 < \tau < 1$
 - 2: **while** $f(\mathbf{x} + \alpha \mathbf{d}) > f(\mathbf{x}) + \gamma_1 \alpha \nabla f(\mathbf{x})^\top \mathbf{d}$ **do**
 - 3: Decrease α : $\alpha \leftarrow \tau \cdot \alpha$
 - 4: **end while**
-




► [Click for source](#)

WOLFE CONDITIONS

- Backtracking is simple and shows good performance in practice
- **But:** Two undesirable scenarios
 - ❶ Initial step size α_{init} is too large $\Rightarrow$ need multiple evaluations of f
 - ❷ Step size is too small with highly negative slopes
- **Solution** for small step sizes:
 - Fix γ_2 with $0 < \gamma_1 < \gamma_2 < 1$.
 - α satisfies **sufficient curvature condition** in $\mathbf{x}$ for $\mathbf{d}$ if

$$|\nabla f(\mathbf{x} + \alpha \mathbf{d})^\top \mathbf{d}| \leq \gamma_2 |\nabla f(\mathbf{x})^\top \mathbf{d}|.$$


Armijo rule + sufficient curvature condition = **Wolfe conditions**




WOLFE CONDITIONS

- **Algorithm** for finding a Wolfe point (point satisfying Wolfe conditions):

- 1 **Bracketing**: Find interval containing Wolfe point
- 2 **Pinpointing**: Find Wolfe point in interval from bracketing




► [Click for source](#)



► [Click for source](#)



Left: Bracketing. Right: Pinpointing.



BRACKETING & PINPOINTING

- **Example:**

- Large initial step size results in quick bracketing but multiple pinpointing steps (**left**).
- Small initial step size results in multiple bracketing steps but quick pinpointing (**right**).



► [Click for source](#)

