



Optimization in Machine Learning

First order methods

Momentum on quadratic forms



Learning goals

- Momentum update in Eigenspace
- Effect of φ

MOMENTUM UPDATE

$$\boldsymbol{\nu}^{[t+1]} = \varphi \boldsymbol{\nu}^{[t]} + \alpha \nabla f(\mathbf{x}^{[t]})$$


$$\mathbf{x}^{[t+1]} = \mathbf{x}^{[t]} - \boldsymbol{\nu}^{[t+1]}$$

which simplifies to

$$\boldsymbol{\nu}^{[t+1]} = \varphi \boldsymbol{\nu}^{[t]} + \alpha (\mathbf{A} \mathbf{x}^{[t]} - \mathbf{b})$$

$$\mathbf{x}^{[t+1]} = \mathbf{x}^{[t]} - \boldsymbol{\nu}^{[t+1]}$$

for the quadratic form.



DYNAMICS OF MOMENTUM

Change basis as before with $\mathbf{w}^{[t]} = \mathbf{V}^\top (\mathbf{x}^{[t]} - \mathbf{x}^*)$ and $\mathbf{u}^{[t]} = \mathbf{V}\boldsymbol{\nu}^{[t]}$, again each component acts independently, but $w_i^{[t]}$ and $u_i^{[t]}$ are coupled:

$$\begin{aligned}u_i^{[t+1]} &= \varphi u_i^{[t]} + \alpha \lambda_i w_i^{[t]} \\w_i^{[t+1]} &= w_i^{[t]} - u_i^{[t+1]}\end{aligned}$$

We rewrite this:

$$\begin{pmatrix} 1 & 0 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} u_i^{[t+1]} \\ w_i^{[t+1]} \end{pmatrix} = \begin{pmatrix} \varphi & \alpha \lambda_i \\ 0 & 1 \end{pmatrix} \begin{pmatrix} u_i^{[t]} \\ w_i^{[t]} \end{pmatrix}$$


inverting the matrix on the LHS, and unravel the recursion:

$$\begin{pmatrix} 1 & 0 \\ 1 & 1 \end{pmatrix}^{-1} = \begin{pmatrix} 1 & 0 \\ -1 & 1 \end{pmatrix} \quad ; \quad \begin{pmatrix} u_i^{[t+1]} \\ w_i^{[t+1]} \end{pmatrix} = \begin{pmatrix} \varphi & \alpha \lambda_i \\ -\varphi & 1 - \alpha \lambda_i \end{pmatrix} \begin{pmatrix} u_i^{[t]} \\ w_i^{[t]} \end{pmatrix} = R^{t+1} \begin{pmatrix} u_i^0 \\ w_i^0 \end{pmatrix}$$

Taking a 2×2 matrix to the t^{th} power can be expressed via its eigenvalues, σ_1 and σ_2 , where $R_j = \frac{R - \sigma_j I}{\sigma_1 - \sigma_2}$:

$$R^t = \begin{cases} \sigma_1^t R_1 - \sigma_2^t R_2, & \text{if } \sigma_1 \neq \sigma_2 \\ \sigma_1^t (tR/\sigma_1 - (t-1)I), & \text{if } \sigma_1 = \sigma_2 \end{cases}$$

- Careful, R is not symmetric, so the EVs can be complex
- In contrast to GD, where we got one geometric series, we have two coupled series with real or complex values



EIGENVALUES OF RECURSION MATRIX


The eigenvalues of an arbitrary 2×2 matrix are:

$$A = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \quad \text{and} \quad \lambda_{1,2} = \frac{\text{tr}(A) \pm \sqrt{\text{tr}(A)^2 - 4 \det(A)}}{2}$$



For us this is:

$$\sigma_{1,2} = \frac{1 + \varphi - \alpha \lambda_i \pm \sqrt{(1 + \varphi - \alpha \lambda_i)^2 - 4\varphi}}{2}$$

- We need both $|\sigma_1|, |\sigma_2| \leq 1$ for convergence
- For the complex case this reduces to $2\sqrt{\varphi}$, which is surprisingly independent of α and λ_i
- For the real case we cannot simplify






MOMENTUM CONVERGENCE REGIONS






Convergence rate is the slowest of $\max\{|\sigma_1|, |\sigma_2|\}$. Each region shows different convergence behavior.

MOMENTUM WITH COMPLEX EIGENVALUES






The eigenvalues of R are complex and we see low frequency ripples.

MOMENTUM WITH POSITIVE EIGENVALUES






Here, both eigenvalues of R are positive with their norm < 1 .
This behavior resembles gradient descent.

ONE-STEP CONVERGENCE






The step size is $\alpha = 1/\lambda_i$ and $\varphi = 0$ - we converge in one step.

OSCILLATING ITERATES



When $\alpha > 1/\lambda_i$, the iterates flip sign every iteration.

DIVERGING ITERATES




If $\max\{|\sigma_1|, |\sigma_2|\} > 1$, the iterates diverge.

CONVERGENCE CONDITIONS

- If we combine all conditions for convergence, we can see:

$$0 < \alpha \lambda_i < 2 + 2\varphi \quad \text{for} \quad 0 \leq \varphi < 1$$

- Comparing this with the results from before ($\varphi = 0$), we see that we gain a stepsize factor of 2 before we diverge!
- Can obtain global convergence rate by optimizing over α and φ
- More involved, see blogpost for details
- We get $\alpha = \left(\frac{2}{\sqrt{\lambda_1} + \sqrt{\lambda_n}} \right)^2$ and $\varphi = \left(\frac{\sqrt{\lambda_n} - \sqrt{\lambda_1}}{\sqrt{\lambda_n} + \sqrt{\lambda_1}} \right)^2$
- Results in convergence rate of $\frac{\sqrt{\kappa}-1}{\sqrt{\kappa}+1}$



PRACTICAL PARAMETER CHOICES

- Compared to GD with $\frac{\kappa-1}{\kappa+1}$ this is much better as the condition is rooted
- Of course, this would in principle require knowledge of the EVs λ_i
- But we can derive simple rule-of-thumb: for poorly conditioned problems, the stepsize is approximately twice that of GD and φ close to 1
- So we want to set φ to a high value and then still pick the highest α which still converges

