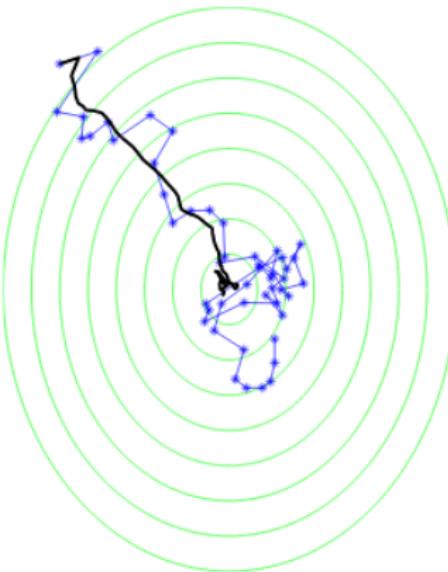


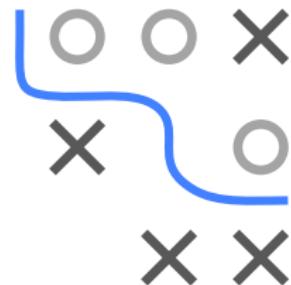
Optimization in Machine Learning

First order methods SGD



Learning goals

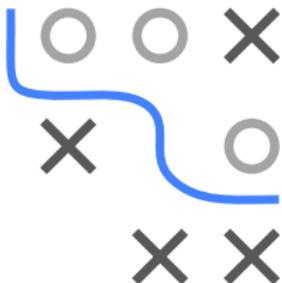
- Understand stochastic gradient descent
- Understand stochasticity and variance of SGD
- Know convergence properties of SGD
- Understand effect of batch size



STOCHASTIC GRADIENT DESCENT

- NB: We use g instead of f as objective, because f is used as model in ML
- $g : \mathbb{R}^d \rightarrow \mathbb{R}$ objective, g is **average over functions**:

$$g(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n g_i(\mathbf{x}), \quad g \text{ and } g_i \text{ smooth}$$



- Stochastic gradient descent (SGD) approximates the gradient

$$\nabla_{\mathbf{x}} g(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\mathbf{x}} g_i(\mathbf{x}) := \mathbf{d} \quad \text{by}$$

$$\frac{1}{|J|} \sum_{i \in J} \nabla_{\mathbf{x}} g_i(\mathbf{x}) := \hat{\mathbf{d}},$$

with random subset $J \subset \{1, 2, \dots, n\}$ of gradients called
mini-batch

- This is done e.g. when computing the true gradient is **expensive**

SGD ALGORITHM

Algorithm Basic SGD pseudo code

```
1: Initialize  $\mathbf{x}^{[0]}$ ,  $t = 0$ 
2: while stopping criterion not met do
3:   Randomly shuffle indices and partition into minibatches  $J_1, \dots, J_K$  of size  $m$ 
4:   for  $k \in \{1, \dots, K\}$  do
5:      $t \leftarrow t + 1$ 
6:     Compute gradient estimate with  $J_k$ :  $\hat{\mathbf{d}}^{[t]} \leftarrow \frac{1}{m} \sum_{i \in J_k} \nabla_{\mathbf{x}} g_i(\mathbf{x}^{[t-1]})$ 
7:     Apply update:  $\mathbf{x}^{[t]} \leftarrow \mathbf{x}^{[t-1]} - \alpha \cdot \hat{\mathbf{d}}^{[t]}$ 
8:   end for
9: end while
```



- Instead of drawing batches randomly we might want to go through the g_i sequentially (unless g_i are sorted in any way)
- Updates are computed faster, but also more stochastic:
 - In the simplest case, batch-size $m := |J_k|$ is set to $m = 1$
 - If n is a billion, computation of update is a billion times faster
 - **But** (later): Convergence rates suffer from stochasticity!

SGD IN ML



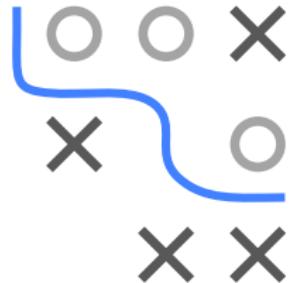
- In ML, we perform ERM:

$$\mathcal{R}(\boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n \underbrace{L\left(y^{(i)}, f(\mathbf{x}^{(i)} | \boldsymbol{\theta})\right)}_{g_i(\boldsymbol{\theta})}$$

- for a data set $\mathcal{D} = ((\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)}))$
- a loss function $L(y, f(\mathbf{x}))$, e.g., L2 loss $L(y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2$
- and a model class f , e.g., the linear model $f(\mathbf{x}^{(i)} | \boldsymbol{\theta}) = \boldsymbol{\theta}^\top \mathbf{x}$

- For large data sets, computing the exact gradient

$$\mathbf{d} = \frac{1}{n} \sum_{i=1}^n \nabla_{\theta} L \left(y^{(i)}, f(\mathbf{x}^{(i)} | \theta) \right)$$



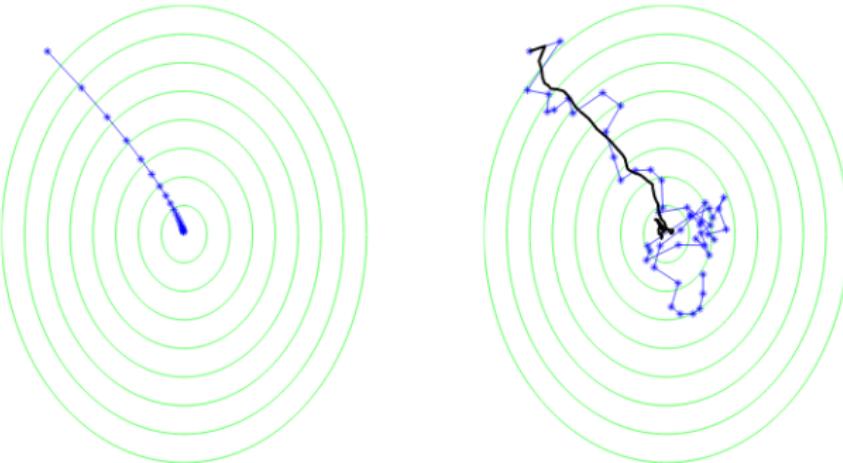
may be expensive or even infeasible to compute and is approximated by

$$\hat{\mathbf{d}} = \frac{1}{m} \sum_{i \in J} \nabla_{\theta} L \left(y^{(i)}, f(\mathbf{x}^{(i)} | \theta) \right),$$

for $J \subset \{1, 2, \dots, n\}$ random subset

- NB:** Often, maximum size of J technically limited by memory size

STOCHASTICITY OF SGD



▶ Click for source

Minimize $g(x_1, x_2) = 1.25(x_1 + 6)^2 + (x_2 - 8)^2$ **Left:** GD, **Right:** SGD
Black line shows average value across multiple runs



STOCHASTICITY OF SGD

- Assume batch size $m = 1$ (statements also apply for larger batches)
- **(Possibly) suboptimal direction:** Approximate gradient $\hat{\mathbf{d}} = \nabla_{\mathbf{x}} g_i(\mathbf{x})$ might point in suboptimal (possibly not even a descent!) direction
- **Unbiased estimate:** If J drawn i.i.d., approximate gradient $\hat{\mathbf{d}}$ is an unbiased estimate of gradient $\mathbf{d} = \nabla_{\mathbf{x}} g(\mathbf{x}) = \sum_{i=1}^n \nabla_{\mathbf{x}} g_i(\mathbf{x})$:

$$\begin{aligned}\mathbb{E}_i [\nabla_{\mathbf{x}} g_i(\mathbf{x})] &= \sum_{i=1}^n \nabla_{\mathbf{x}} g_i(\mathbf{x}) \cdot \mathbb{P}(i = i) \\ &= \sum_{i=1}^n \nabla_{\mathbf{x}} g_i(\mathbf{x}) \cdot \frac{1}{n} = \nabla_{\mathbf{x}} g(\mathbf{x}).\end{aligned}$$

- **Conclusion:** SGD might perform single suboptimal moves, but moves in “right direction” **on average**

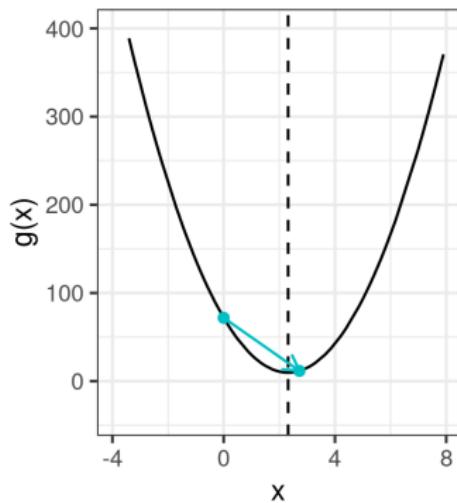


ERRATIC BEHAVIOR OF SGD

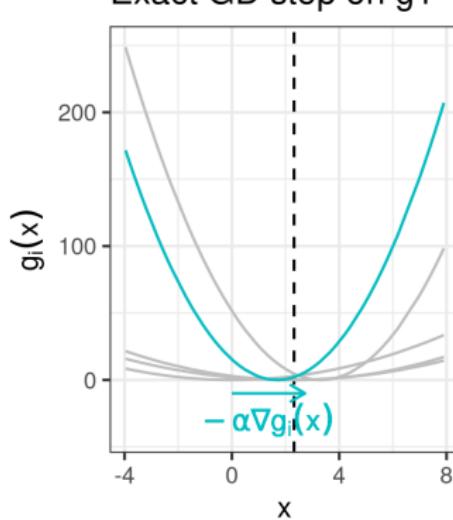
Example: $g(\mathbf{x}) = \sum_{i=1}^5 g_i(\mathbf{x})$, g_i quadratic, batch size $m = 1$



SGD step 1 on $g(\mathbf{x})$



Exact GD step on g_1

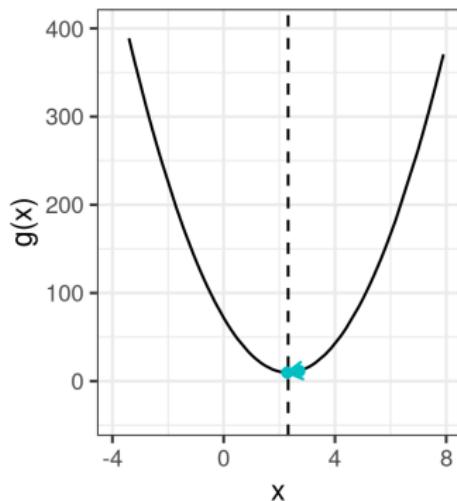


ERRATIC BEHAVIOR OF SGD

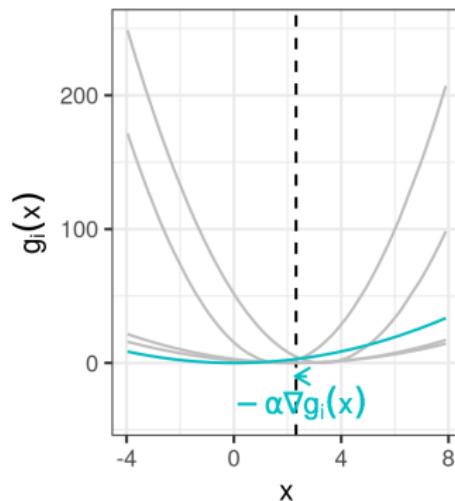
Example: $g(\mathbf{x}) = \sum_{i=1}^5 g_i(\mathbf{x})$, g_i quadratic, batch size $m = 1$



SGD step 2 on $g(\mathbf{x})$



Exact GD step on g_2

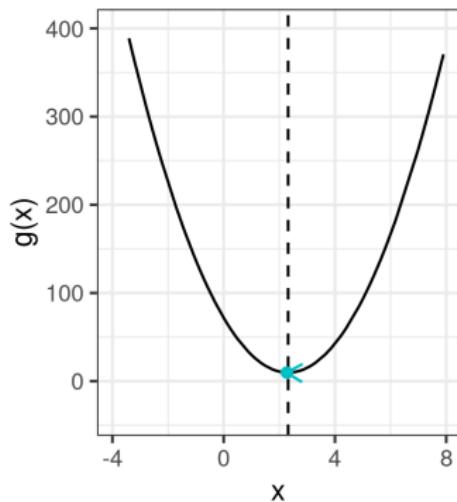


ERRATIC BEHAVIOR OF SGD

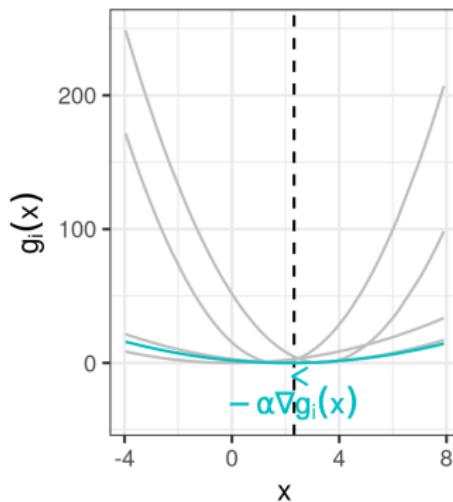
Example: $g(\mathbf{x}) = \sum_{i=1}^5 g_i(\mathbf{x})$, g_i quadratic, batch size $m = 1$



SGD step 3 on $g(\mathbf{x})$



Exact GD step on g_3

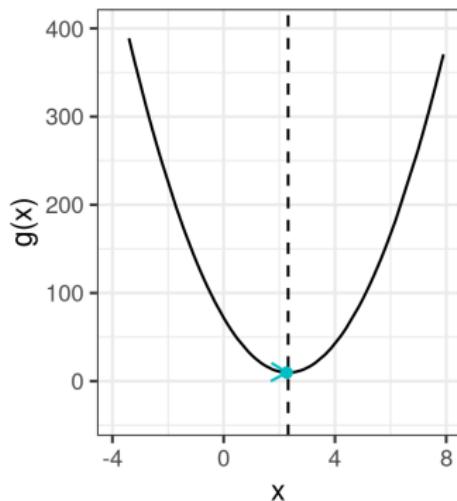


ERRATIC BEHAVIOR OF SGD

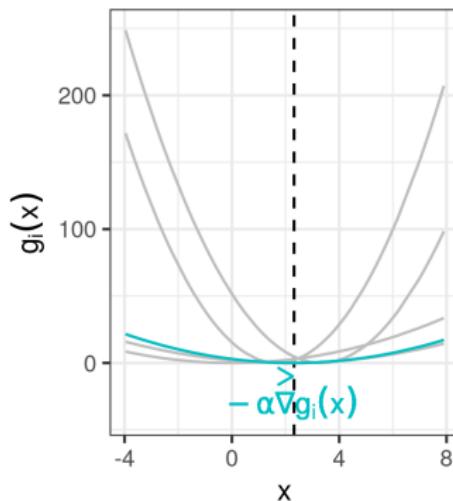
Example: $g(\mathbf{x}) = \sum_{i=1}^5 g_i(\mathbf{x})$, g_i quadratic, batch size $m = 1$



SGD step 4 on $g(\mathbf{x})$



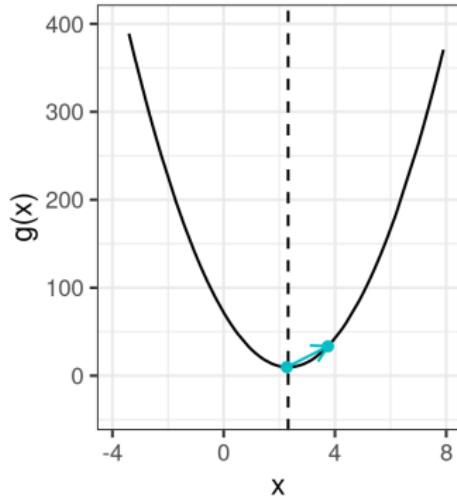
Exact GD step on g_i



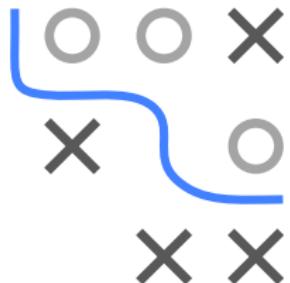
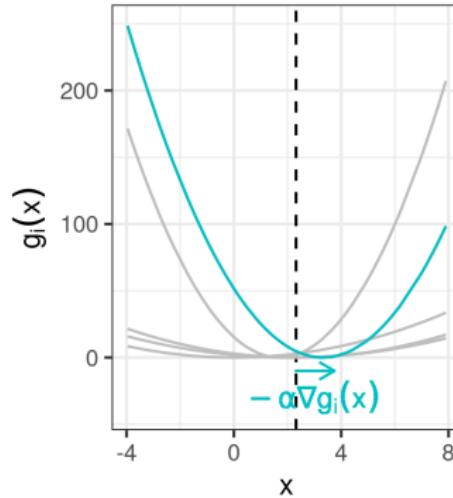
ERRATIC BEHAVIOR OF SGD

Example: $g(\mathbf{x}) = \sum_{i=1}^5 g_i(\mathbf{x})$, g_i quadratic, batch size $m = 1$

SGD step 5 on $g(\mathbf{x})$

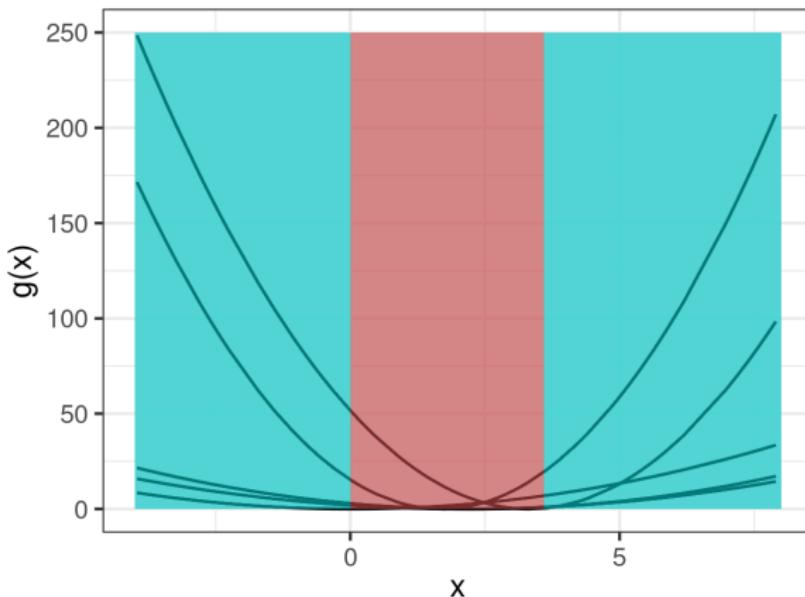


Exact GD step on g_5



In iteration 5, SGD performs a suboptimal move away from the minimum.

ERRATIC BEHAVIOR OF SGD



Blue area: Each $-\nabla g_i(\mathbf{x})$ points towards minimum

Red area ("confusion area"): $-\nabla g_i(\mathbf{x})$ might point away from minimum
and perform a suboptimal move

ERRATIC BEHAVIOR OF SGD

- At location \mathbf{x} , “confusion” is captured by variance of gradients

$$\frac{1}{n} \sum_{i=1}^n \|\nabla_{\mathbf{x}} g_i(\mathbf{x}) - \nabla_{\mathbf{x}} g(\mathbf{x})\|^2$$

- If term is 0, next step goes in gradient direction (for each i)
- If term is small, next step *likely* goes in gradient direction
- If term is large, next step likely goes in direction different than gradient



CONVERGENCE OF SGD

- As a consequence, SGD has worse convergence properties than GD
- But:** Can be controlled via **increasing batches** or **reducing step size**

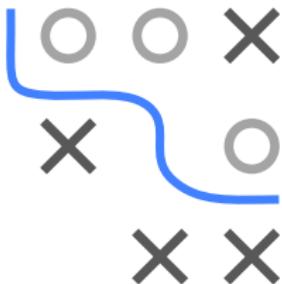
The larger the batch size m

- the better the approximation to $\nabla_{\mathbf{x}} g(\mathbf{x})$
- the lower the variance
- the lower the risk of performing steps in the wrong direction

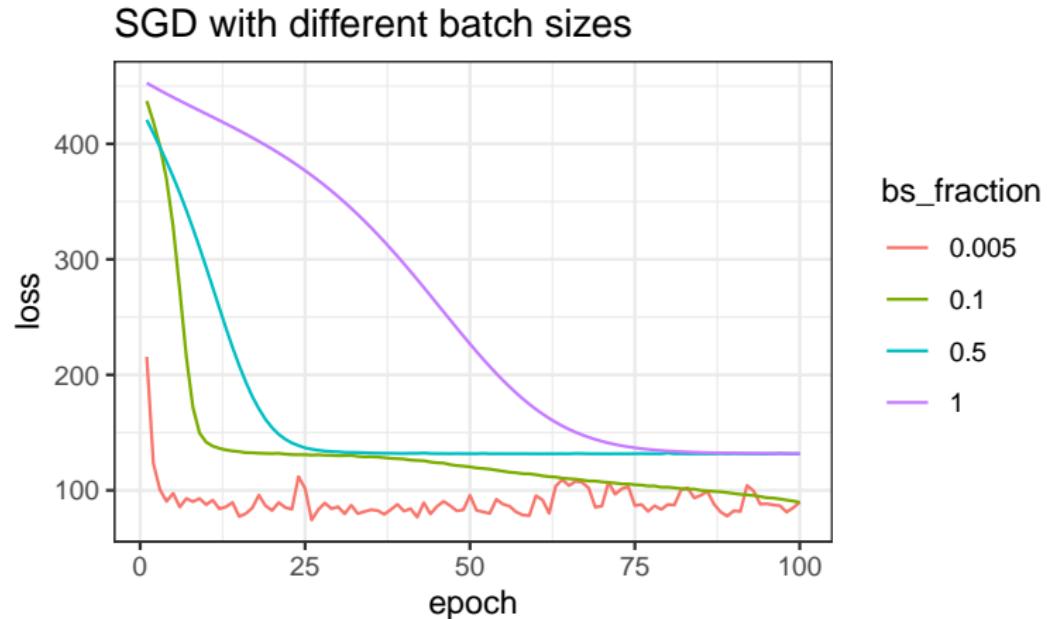
The smaller the step size α

- the smaller a step in a potentially wrong direction
- the lower the effect of high variance

As maximum batch size is usually limited by computational resources (memory), choosing the step size is crucial



EFFECT OF BATCH SIZE



bs_fraction

- 0.005
- 0.1
- 0.5
- 1



SGD for a NN with batch size $\in \{0.5\%, 10\%, 50\%\}$ of the training data
The higher the batch size, the lower the variance