


Optimization in Machine Learning

First order methods SGD Further Details



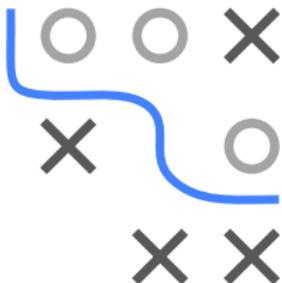

Learning goals

- Decreasing step size for SGD
- Stopping rules
- SGD with momentum



SGD WITH CONSTANT STEP SIZE

- Example: SGD with constant step size




- Fast convergence of SGD initially
- Erratic behavior later (variance too big)

SGD WITH DECREASING STEP SIZE

- Idea: Decrease step size to reduce magnitude of erratic steps
- Trade-off:
 - If step size $\alpha^{[t]}$ decreases slowly, large erratic steps
 - If step size decreases too fast, performance is impaired
- SGD converges for sufficiently smooth functions if

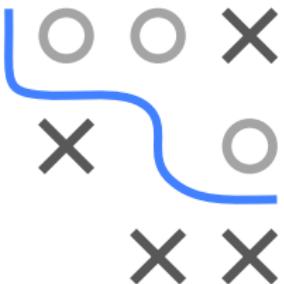

$$\frac{\sum_{t=1}^{\infty} (\alpha^{[t]})^2}{\sum_{t=1}^{\infty} \alpha^{[t]}} = 0$$

(“how much noise affects you” by “how far you can get”)



SGD WITH DECREASING STEP SIZE

- Popular solution: step size fulfilling $\alpha^{[t]} \in \mathcal{O}(1/t)$



- Example continued: Step size $\alpha^{[t]} = 0.2/t$
- Often not working well in practice: step size gets small quite fast
- Alternative: $\alpha^{[t]} \in \mathcal{O}(1/\sqrt{t})$

ADVANCED STEP SIZE CONTROL

Why not Armijo-based step size control?

- Backtracking line search or other approaches based on Armijo rule usually not suitable: Armijo condition

$$g(\mathbf{x} + \alpha \mathbf{d}) \leq g(\mathbf{x}) + \gamma_1 \alpha \nabla g(\mathbf{x})^\top \mathbf{d}$$

requires evaluating full gradient

- But SGD is used to *avoid* expensive gradient computations
- Research aims at finding inexact line search methods that provide better convergence behaviour, e.g., Vaswani et al., *Painless Stochastic Gradient: Interpolation, Line-Search, and Convergence Rates*. NeurIPS, 2019




MINI-BATCHES

- Reduce noise by increasing batch size m for better approximation


$$\hat{\mathbf{d}} = \frac{1}{m} \sum_{i \in J} \nabla_{\mathbf{x}} g_i(\mathbf{x}) \approx \frac{1}{n} \sum_{i=1}^n \nabla_{\mathbf{x}} g_i(\mathbf{x}) = \mathbf{d}$$

- Usually, the batch size is limited by computational resources (e.g., how much data you can load into the memory)



- Example continued: Batch size $m = 1$ vs. $m = 5$

STOPPING RULES FOR SGD



- For GD: We usually stop when gradient is close to 0 (i.e., we are close to a stationary point)
- For SGD: Individual gradients do not necessarily go to zero, and we cannot access full gradient
- Practicable solution for ML:
 - Measure the validation set error after T iterations
 - Stop if validation set error is not improving

SGD AND ML


General remarks:

- SGD is a variant of GD
- SGD particularly suitable for large-scale ML when evaluating gradient is too expensive / restricted by computational resources
- SGD and variants are the most commonly used methods in modern ML, for example:
 - Linear models
Note that even for the linear model and quadratic loss, where a closed form solution is available, SGD might be used if the size n of the dataset is too large and the design matrix does not fit into memory
 - Neural networks
 - Support vector machines
 - ...



SGD WITH MOMENTUM

- SGD is usually used with momentum due to reasons mentioned in previous chapters



Algorithm Stochastic gradient descent with momentum

- 1: **require** step size α and momentum φ
- 2: **require** initial parameter \mathbf{x} and initial velocity \mathbf{v}
- 3: **while** stopping criterion not met **do**
- 4: Sample mini-batch of m examples
- 5: Compute gradient estimate $\nabla \hat{g}(\mathbf{x})$ using mini-batch
- 6: Compute velocity update: $\mathbf{v} \leftarrow \varphi \mathbf{v} - \alpha \nabla \hat{g}(\mathbf{x})$
- 7: Apply update: $\mathbf{x} \leftarrow \mathbf{x} + \mathbf{v}$
- 8: **end while**
