


Optimization in Machine Learning

First order methods Adam and friends





Learning goals

- Adaptive step sizes
- AdaGrad
- RMSProp
- Adam



ADAPTIVE STEP SIZES

- Step size is probably the most important control parameter
- Has strong influence on performance
- Natural to use different step size for each input individually and automatically adapt them



ADAGRAD

- AdaGrad adapts step sizes by scaling them inversely proportional to square root of the sum of the past squared derivatives
 - Inputs with large derivatives get smaller step sizes
 - Inputs with small derivatives get larger step sizes
- Accumulation of squared gradients can result in premature small step sizes
(Goodfellow et al., 2016)



Algorithm AdaGrad

```
1: require Global step size  $\alpha$ 
2: require Initial parameter  $\theta$ 
3: require Small constant  $\beta$ , perhaps  $10^{-7}$ , for numerical stability
4: Initialize gradient accumulation variable  $r = \mathbf{0}$ 
5: while stopping criterion not met do
6:   Sample minibatch of  $m$  examples from the training set  $\{\tilde{x}^{(1)}, \dots, \tilde{x}^{(m)}\}$ 
7:   Compute gradient estimate:  $\hat{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(y^{(i)}, f(\tilde{x}^{(i)} | \theta))$ 
8:   Accumulate squared gradient  $r \leftarrow r + \hat{g} \odot \hat{g}$  where  $\odot$ : element-wise product
9:   Compute update:  $\nabla_{\theta} = -\frac{\alpha}{\beta + \sqrt{r}} \odot \hat{g}$ 
10:  Apply update:  $\theta \leftarrow \theta + \nabla_{\theta}$ 
11: end while
```

RMSProp

- Modification of AdaGrad
- Resolves AdaGrad's radically diminishing step sizes
- Gradient accumulation is replaced by exponentially weighted moving average
- Theoretically, leads to performance gains in non-convex scenarios
- Empirically, RMSProp is a very effective optimization algorithm. Particularly, it is employed routinely by DL practitioners




Algorithm RMSProp

- 1: **require** Global step size α and decay rate $\rho \in [0, 1)$
 - 2: **require** Initial parameter θ
 - 3: **require** Small constant β , perhaps 10^{-6} , for numerical stability
 - 4: Initialize gradient accumulation variable $r = \mathbf{0}$
 - 5: **while** stopping criterion not met **do**
 - 6: Sample minibatch of m examples from the training set $\{\tilde{x}^{(1)}, \dots, \tilde{x}^{(m)}\}$
 - 7: Compute gradient estimate: $\hat{g} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(y^{(i)}, f(\tilde{x}^{(i)} | \theta))$
 - 8: Accumulate squared gradient $r \leftarrow \rho r + (1 - \rho)\hat{g} \odot \hat{g}$
 - 9: Compute update: $\nabla_{\theta} \theta = -\frac{\alpha}{\beta + \sqrt{r}} \odot \hat{g}$
 - 10: Apply update: $\theta \leftarrow \theta + \nabla_{\theta} \theta$
 - 11: **end while**
-

ADAM

- Adaptive Moment Estimation also has adaptive step sizes
- Uses the 1st and 2nd moments of gradients
 - Keeps an exponentially decaying average of past gradients (1st moment)
 - Like RMSProp, stores an exp-decaying avg of past squared gradients (2nd moment)
 - Can be seen as combo of RMSProp + momentum



ADAM

Algorithm Adam

- 1: **require** Global step size α (suggested default: 0.001)
- 2: **require** Exponential decay rates for moment estimates, ρ_1 and ρ_2 in $[0, 1]$ (suggested defaults: 0.9 and 0.999 respectively)
- 3: **require** Small constant β (suggested default 10^{-8})
- 4: **require** Initial parameters θ
- 5: Initialize time step $t = 0$
- 6: Initialize 1st and 2nd moment variables $s^{[0]} = 0, r^{[0]} = 0$
- 7: **while** stopping criterion not met **do**
- 8: $t \leftarrow t + 1$
- 9: Sample a minibatch of m examples from the training set $\{\tilde{x}^{(1)}, \dots, \tilde{x}^{(m)}\}$
- 10: Compute gradient estimate: $\hat{g}^{[t]} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(y^{(i)}, f(\tilde{x}^{(i)} | \theta))$
- 11: Update biased first moment estimate: $s^{[t]} \leftarrow \rho_1 s^{[t-1]} + (1 - \rho_1) \hat{g}^{[t]}$
- 12: Update biased second moment estimate: $r^{[t]} \leftarrow \rho_2 r^{[t-1]} + (1 - \rho_2) \hat{g}^{[t]} \odot \hat{g}^{[t]}$
- 13: Correct bias in first moment: $\hat{s} \leftarrow \frac{s^{[t]}}{1 - \rho_1^t}$
- 14: Correct bias in second moment: $\hat{r} \leftarrow \frac{r^{[t]}}{1 - \rho_2^t}$
- 15: Compute update: $\nabla \theta = -\alpha \frac{\hat{s}}{\sqrt{\hat{r}} + \beta}$
- 16: Apply update: $\theta \leftarrow \theta + \nabla \theta$
- 17: **end while**



ADAM: BIAS IN MOMENT ESTIMATES

- Initializes moment variables \mathbf{s} and \mathbf{r} with zero \Rightarrow Bias towards zero

$$\mathbb{E}[\mathbf{s}^{[t]}] \neq \mathbb{E}[\hat{\mathbf{g}}^{[t]}] \quad \text{and} \quad \mathbb{E}[\mathbf{r}^{[t]}] \neq \mathbb{E}[\hat{\mathbf{g}}^{[t]} \odot \hat{\mathbf{g}}^{[t]}]$$

(\mathbb{E} calculated over minibatches)

- Indeed: Unrolling $\mathbf{s}^{[t]}$ yields

$$\mathbf{s}^{[0]} = 0$$

$$\mathbf{s}^{[1]} = \rho_1 \mathbf{s}^{[0]} + (1 - \rho_1) \hat{\mathbf{g}}^{[1]} = (1 - \rho_1) \hat{\mathbf{g}}^{[1]}$$

$$\mathbf{s}^{[2]} = \rho_1 \mathbf{s}^{[1]} + (1 - \rho_1) \hat{\mathbf{g}}^{[2]} = \rho_1 (1 - \rho_1) \hat{\mathbf{g}}^{[1]} + (1 - \rho_1) \hat{\mathbf{g}}^{[2]}$$

$$\mathbf{s}^{[3]} = \rho_1 \mathbf{s}^{[2]} + (1 - \rho_1) \hat{\mathbf{g}}^{[3]} = \rho_1^2 (1 - \rho_1) \hat{\mathbf{g}}^{[1]} + \rho_1 (1 - \rho_1) \hat{\mathbf{g}}^{[2]} + (1 - \rho_1) \hat{\mathbf{g}}^{[3]}$$

- Therefore: $\mathbf{s}^{[t]} = (1 - \rho_1) \sum_{i=1}^t \rho_1^{t-i} \hat{\mathbf{g}}^{[i]}$


- Note:** Contributions of past $\hat{\mathbf{g}}^{[i]}$ decreases rapidly



ADAM: BIAS CORRECTION

- We continue with


$$\begin{aligned}\mathbb{E}[\mathbf{s}^{[t]}] &= \mathbb{E}[(1 - \rho_1) \sum_{i=1}^t \rho_1^{t-i} \hat{\mathbf{g}}^{[i]}] \\ &= \mathbb{E}[\hat{\mathbf{g}}^{[t]}](1 - \rho_1) \sum_{i=1}^t \rho_1^{t-i} + \zeta \\ &= \mathbb{E}[\hat{\mathbf{g}}^{[t]}](1 - \rho_1^t) + \zeta,\end{aligned}$$




where we approximated $\hat{\mathbf{g}}^{[i]}$ by $\hat{\mathbf{g}}^{[t]}$. The resulting error is put in ζ and be kept small due to the exponential weights of past gradients

- Therefore: $\mathbf{s}^{[t]}$ is a biased estimator of $\hat{\mathbf{g}}^{[t]}$
- But bias vanishes for $t \rightarrow \infty$ ($\rho_1^t \rightarrow 0$)
- Ignoring ζ , we correct for the bias by $\hat{\mathbf{s}}^{[t]} = \frac{\mathbf{s}^{[t]}}{(1 - \rho_1^t)}$
- Analogously: $\hat{\mathbf{r}}^{[t]} = \frac{\mathbf{r}^{[t]}}{(1 - \rho_2^t)}$

COMPARISON OF OPTIMIZERS: ANIMATION



▶ Click for source



▶ Click for source

Comparison of SGD optimizers near saddle point

Left: After start. **Right:** Later



All methods accelerate compared to vanilla SGD

Best is RMSProp, then AdaGrad (Adam is missing here)

Credits: Dettmers (2015) and Radford



COMPARISON ON QUADRATIC FORM



SGD vs SGD with Momentum vs Adam on a quadratic form