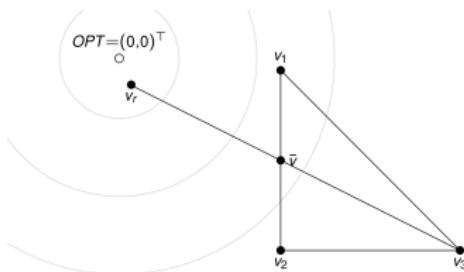


# Optimization in Machine Learning

## Nelder-Mead method



### Learning goals

- General idea
- Reflection, expansion, contraction
- Advantages & disadvantages
- Examples



# NELDER-MEAD METHOD

- Derivative-free method  $\Rightarrow$  heuristic
- Generalization of bisection in  $d$ -dimensional space
- Based on  $d$ -simplex, defined by  $d + 1$  points:
  - $d = 1$  interval
  - $d = 2$  triangle
  - $d = 3$  tetrahedron
  - $\dots$



# NELDER-MEAD METHOD

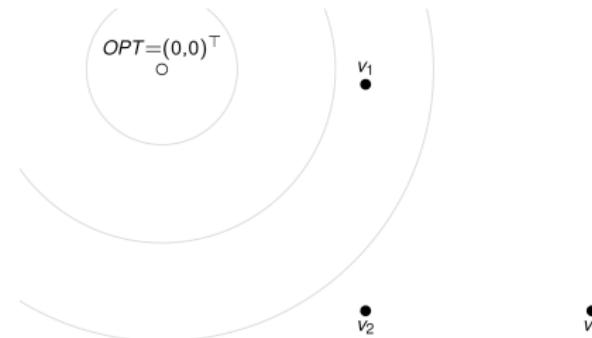
A version of the Nelder-Mead method:

Initialization: choose  $d + 1$  random, affinely independent points  $\mathbf{v}_i$   
( $\mathbf{v}_i$  are vertices: corner points of the simplex/polytope)

1. Order points according to ascending function values

$$f(\mathbf{v}_1) \leq f(\mathbf{v}_2) \leq \dots \leq f(\mathbf{v}_d) \leq f(\mathbf{v}_{d+1})$$

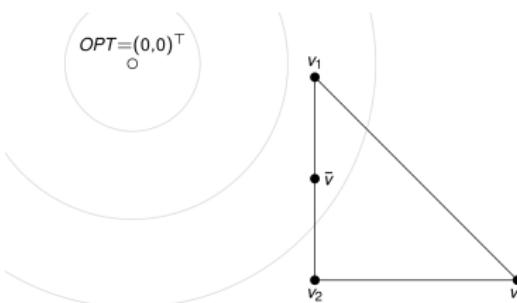
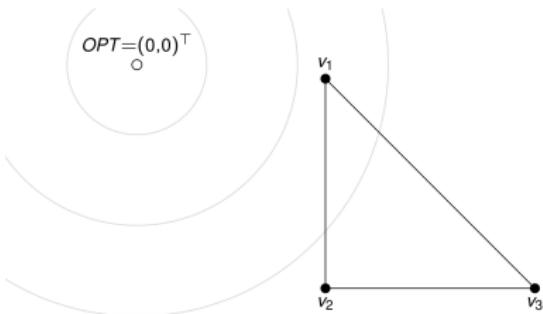
with  $\mathbf{v}_1$  best point,  $\mathbf{v}_{d+1}$  worst point



# NELDER-MEAD METHOD

2. Compute centroid without worst point

$$\bar{\mathbf{v}} = \frac{1}{d} \sum_{i=1}^d \mathbf{v}_i.$$

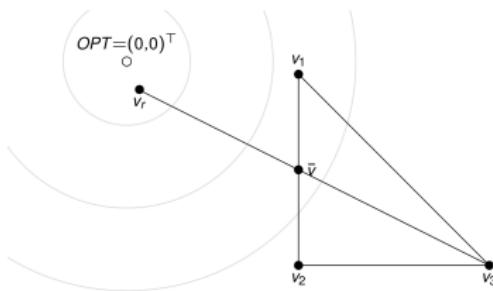


# NELDER-MEAD METHOD

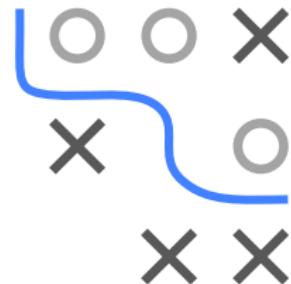
3. Reflection: compute reflection point

$$\mathbf{v}_r = \bar{\mathbf{v}} + \rho(\bar{\mathbf{v}} - \mathbf{v}_{d+1}),$$

with  $\rho > 0$ . Compute  $f(\mathbf{v}_r)$ .



Note: Default value for reflection coefficient:  $\rho = 1$



# NELDER-MEAD METHOD

Distinguish three cases:

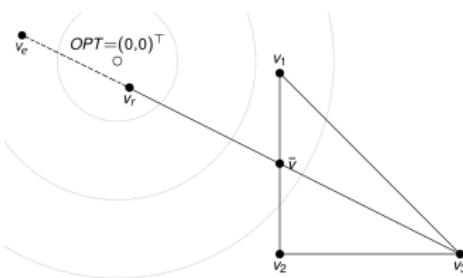
- Case 1:  $f(\mathbf{v}_1) \leq f(\mathbf{v}_r) < f(\mathbf{v}_d) \Rightarrow$  Accept  $\mathbf{v}_r$  and discard  $\mathbf{v}_{d+1}$

- Case 2:  $f(\mathbf{v}_r) < f(\mathbf{v}_1) \Rightarrow$  Expansion:

$$\mathbf{v}_e = \bar{\mathbf{v}} + \chi(\mathbf{v}_r - \bar{\mathbf{v}}), \quad \chi > 1$$

We discard  $\mathbf{v}_{d+1}$  and accept the better of  $\mathbf{v}_r$  and  $\mathbf{v}_e$

Note: default value for expansion coefficient:  $\chi = 2$



# NELDER-MEAD METHOD

- Case 3:  $f(\mathbf{v}_r) \geq f(\mathbf{v}_d) \Rightarrow$  Contraction:

$$\mathbf{v}_c = \bar{\mathbf{v}} + \gamma(\mathbf{v}_{d+1} - \bar{\mathbf{v}})$$

with  $0 < \gamma \leq 1/2$

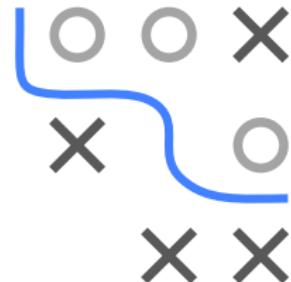
- If  $f(\mathbf{v}_c) < f(\mathbf{v}_{d+1})$ , accept  $\mathbf{v}_c$
- Otherwise, shrink entire simplex (Shrinking):

$$\mathbf{v}_i = \mathbf{v}_1 + \sigma(\mathbf{v}_i - \mathbf{v}_1) \quad \forall i$$

Note: default values for contraction and shrinking coefficients

$$\gamma = \sigma = 1/2$$

- Repeat all steps until stopping criterion met



# NELDER-MEAD SUMMARY

Advantages:

- No gradients needed
- Robust, often works well for non-differentiable functions

Drawbacks:

- Relatively slow (not applicable in high dimensions)
- Not each step improves, only mean of corner values is reduced
- No guarantee for convergence to local optimum / stationary point

Visualization:

<http://www.benfrederickson.com/numerical-optimization/>

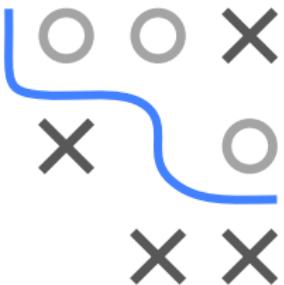
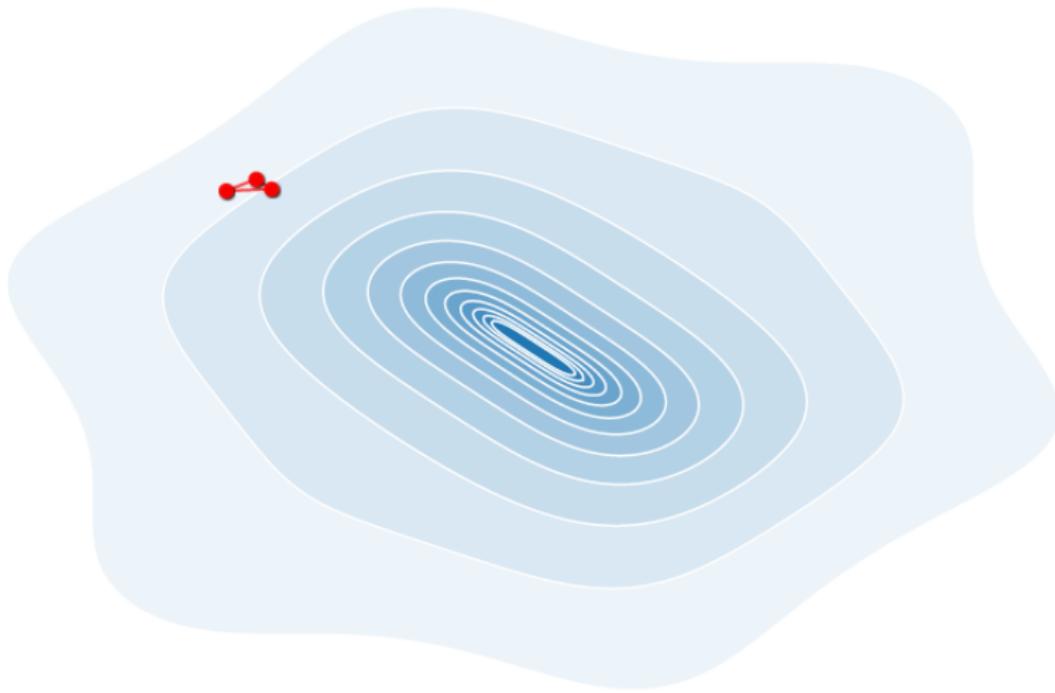
Note: Nelder-Mead is default method of R function `optim()`

If gradient is available and cheap, L-BFGS is preferred



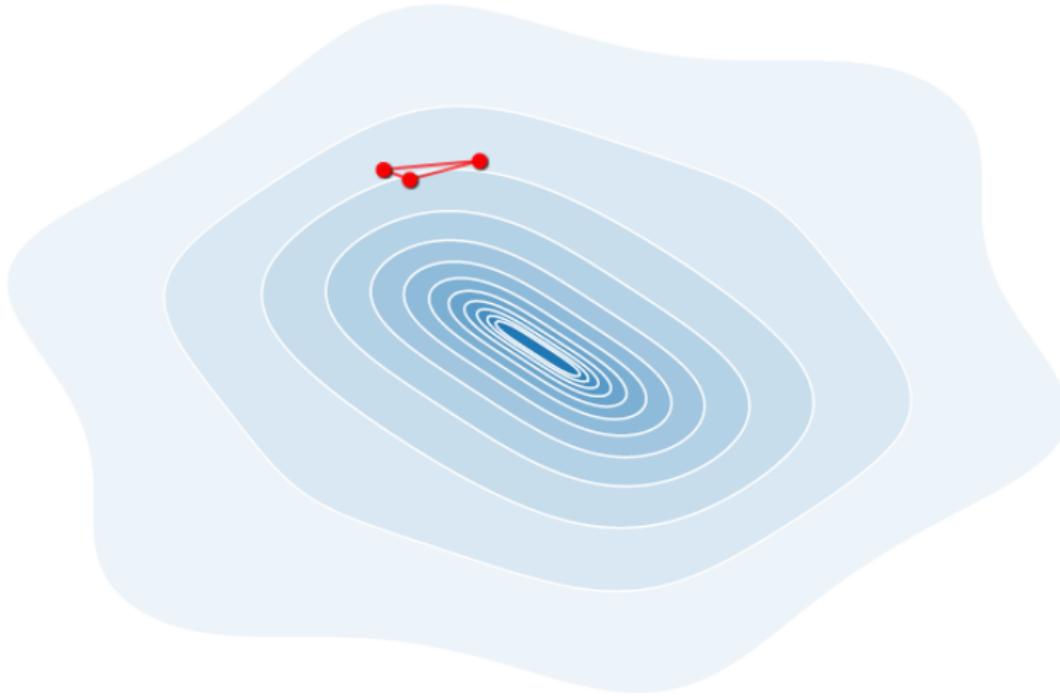
# NELDER-MEAD VISUALIZATION IN 2D

$$\min_{\mathbf{x}} f(\mathbf{x}_1, \mathbf{x}_2) = \mathbf{x}_1^2 + \mathbf{x}_2^2 + \mathbf{x}_1 \cdot \sin \mathbf{x}_2 + \mathbf{x}_2 \cdot \sin \mathbf{x}_1$$



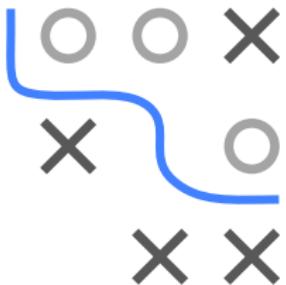
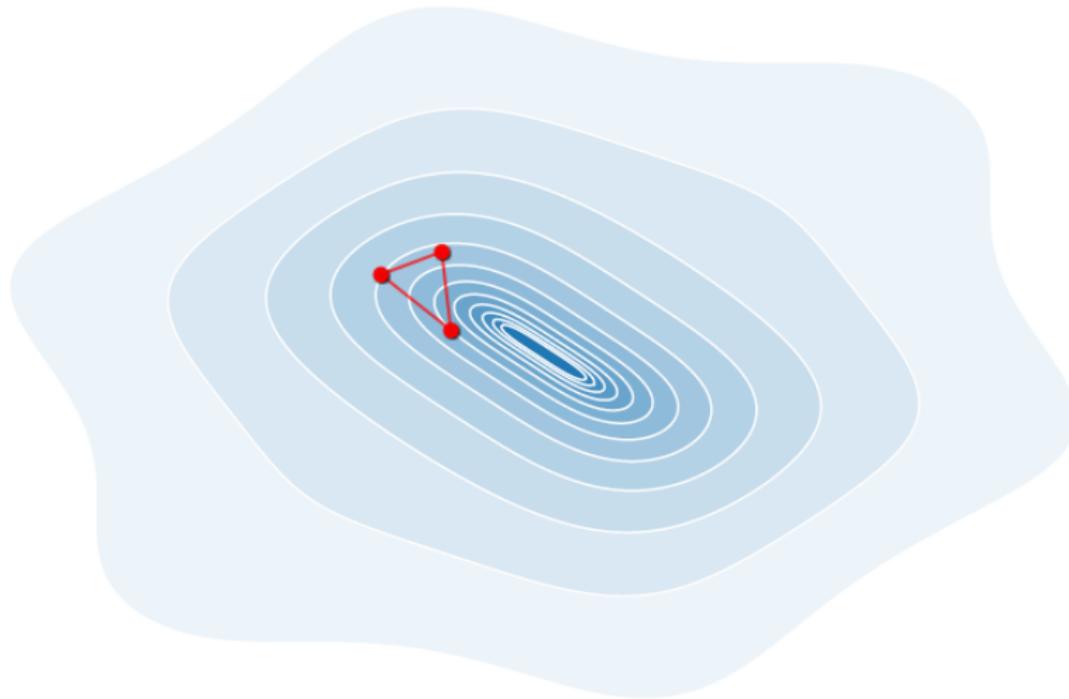
# NELDER-MEAD VISUALIZATION IN 2D

$$\min_{\mathbf{x}} f(x_1, x_2) = x_1^2 + x_2^2 + x_1 \cdot \sin x_2 + x_2 \cdot \sin x_1$$



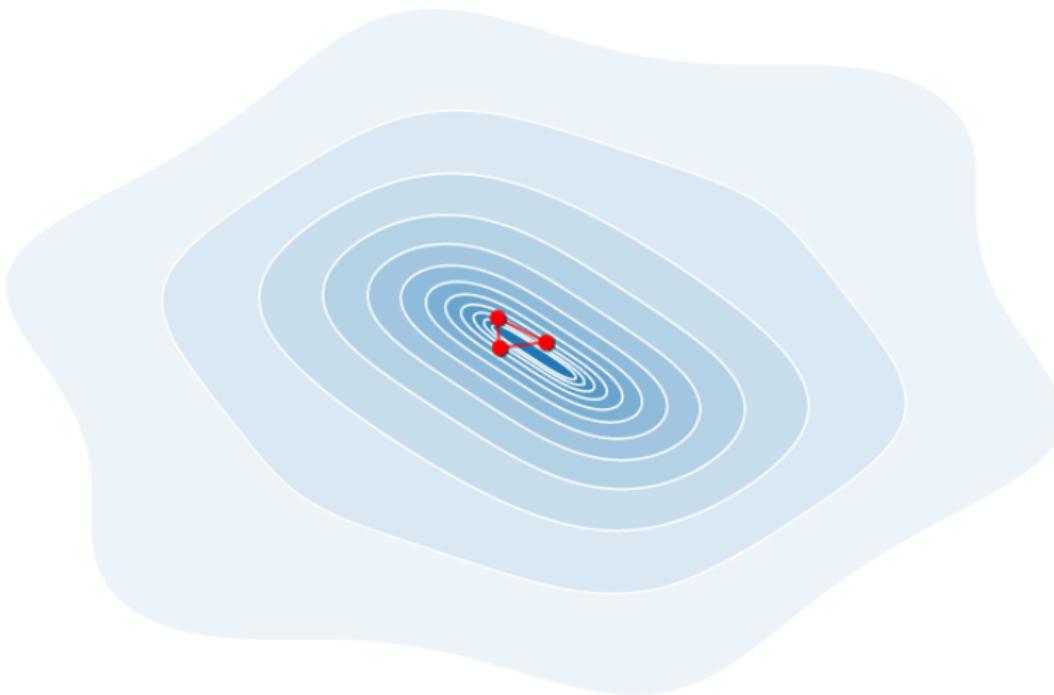
# NELDER-MEAD VISUALIZATION IN 2D

$$\min_{\mathbf{x}} f(x_1, x_2) = x_1^2 + x_2^2 + x_1 \cdot \sin x_2 + x_2 \cdot \sin x_1$$

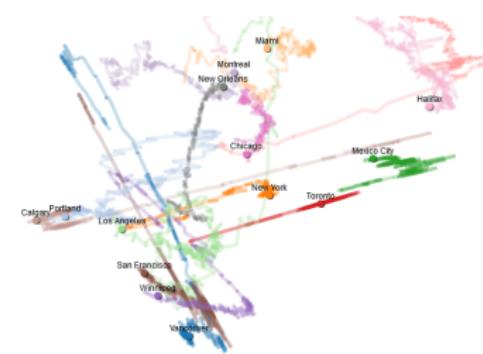
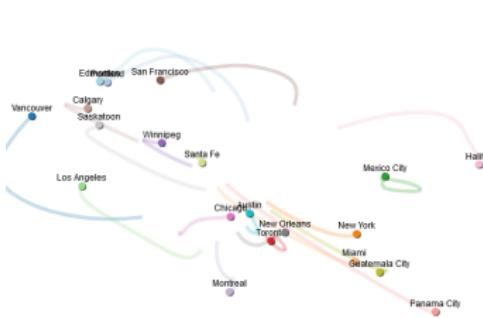


# NELDER-MEAD VISUALIZATION IN 2D

$$\min_{\mathbf{x}} f(x_1, x_2) = x_1^2 + x_2^2 + x_1 \cdot \sin x_2 + x_2 \cdot \sin x_1$$

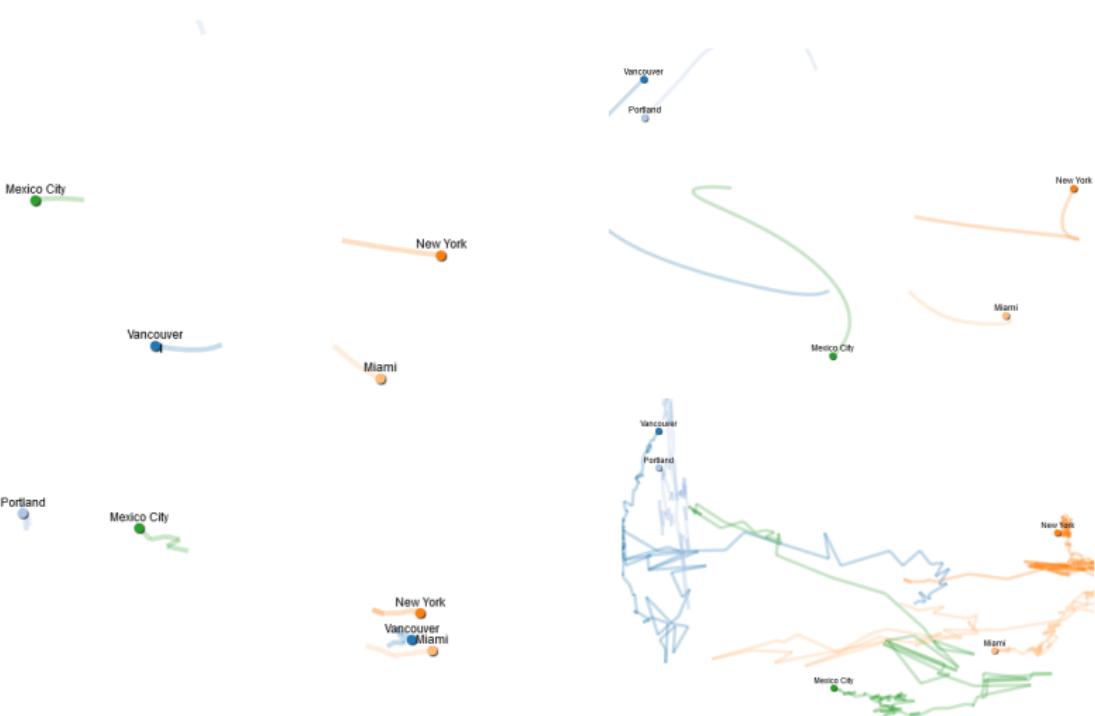


# NELDER-MEAD VS. GD



Nelder-Mead in multiple dimensions: Organize points (US cities) to keep predefined mutual distances. For 10 cities, gradient descent (top) converges well for a suitable learning rate. Nelder-Mead (bottom) fails to converge, even after many iterations.

# NELDER-MEAD VS. GD



Even for only 5 cities, Nelder-Mead (bottom) performs poorly

However, gradient descent (top) still works

