



Optimization in Machine Learning

Nelder-Mead method



Learning goals

- General idea
- Reflection, expansion, contraction
- Advantages & disadvantages
- Examples



NELDER-MEAD METHOD


A version of the Nelder-Mead method:

Initialization: choose $d + 1$ random, affinely independent points $\mathbf{v}_i$
($\mathbf{v}_i$ are vertices: corner points of the simplex/polytope)


1. Order points according to ascending function values

$$f(\mathbf{v}_1) \leq f(\mathbf{v}_2) \leq \dots \leq f(\mathbf{v}_d) \leq f(\mathbf{v}_{d+1})$$

with $\mathbf{v}_1$ best point, $\mathbf{v}_{d+1}$ worst point




Example: Nelder-Mead in $d = 2$ dimensions with $d + 1 = 3$ initialized points




NELDER-MEAD METHOD

2. Compute centroid without worst point

$$\bar{\mathbf{v}} = \frac{1}{d} \sum_{i=1}^d \mathbf{v}_i.$$




Simplex with computed centroid $\bar{\mathbf{v}}$



NELDER-MEAD METHOD

Distinguish three cases:

- Case 1: $f(\mathbf{v}_1) \leq f(\mathbf{v}_r) < f(\mathbf{v}_d) \Rightarrow$
Accept $\mathbf{v}_r$ and discard $\mathbf{v}_{d+1}$




Hypothetical visualization of case
1: acceptance of $\mathbf{v}_r$


- Case 2: $f(\mathbf{v}_r) < f(\mathbf{v}_1) \Rightarrow$
Expansion:

$$\mathbf{v}_e = \bar{\mathbf{v}} + \chi(\mathbf{v}_r - \bar{\mathbf{v}}), \quad \chi > 1$$

We discard $\mathbf{v}_{d+1}$ and accept the
better of $\mathbf{v}_r$ and $\mathbf{v}_e$ (default value
for expansion coefficient: $\chi = 2$)



Visualization of case 2: expansion
with $\chi = 2$, but $\mathbf{v}_r$ is kept since
 $f(\mathbf{v}_e) > f(\mathbf{v}_r)$ (cf. figure above)





NELDER-MEAD METHOD

- Case 3 (continued):
 - If $f(\mathbf{v}_c) \geq f(\mathbf{v}_{d+1})$: shrink entire simplex towards $\mathbf{v}_1$
(Shrinking)

$$\mathbf{v}_i = \mathbf{v}_1 + \sigma(\mathbf{v}_i - \mathbf{v}_1) \quad \forall i$$

Note: default value for shrinking coefficient: $\sigma = 1/2$



Hypothetical visualization of case 3: shrinkage with $\sigma = 1/2$ (if $f(\mathbf{v}_c) \geq f(\mathbf{v}_3)$)

4. Repeat all steps until stopping criterion met

NELDER-MEAD SUMMARY

Advantages:

- No gradients needed
- Robust, often works well for non-differentiable functions

Drawbacks:


- Relatively slow (not applicable in high dimensions)
- Not each step improves, only mean of corner values is reduced
- No guarantee for convergence to local optimum / stationary point

Visualization:

<http://www.benfrederickson.com/numerical-optimization/>



Note: Nelder-Mead is default method of R function `optim()`

If gradient is available and cheap, L-BFGS is preferred





NELDER-MEAD VISUALIZATION IN 2D

$$\min_{\mathbf{x}} f(x_1, x_2) = x_1^2 + x_2^2 + x_1 \cdot \sin x_2 + x_2 \cdot \sin x_1$$





NELDER-MEAD VISUALIZATION IN 2D

$$\min_{\mathbf{x}} f(x_1, x_2) = x_1^2 + x_2^2 + x_1 \cdot \sin x_2 + x_2 \cdot \sin x_1$$





NELDER-MEAD VISUALIZATION IN 2D

$$\min_{\mathbf{x}} f(x_1, x_2) = x_1^2 + x_2^2 + x_1 \cdot \sin x_2 + x_2 \cdot \sin x_1$$







NELDER-MEAD VISUALIZATION IN 2D

$$\min_{\mathbf{x}} f(x_1, x_2) = x_1^2 + x_2^2 + x_1 \cdot \sin x_2 + x_2 \cdot \sin x_1$$





NELDER-MEAD VS. GD



Nelder-Mead in multiple dimensions: Organize points (US cities) to keep predefined mutual distances. For 10 cities, gradient descent (top) converges well for a suitable learning rate. Nelder-Mead (bottom) fails to converge, even after many iterations.

NELDER-MEAD VS. GD



Even for only 5 cities, Nelder-Mead (bottom) performs poorly
However, gradient descent (top) still works